

Todo list

add more about the general setting (time series of data, consider medical cases as examples, ...)	2
add reference	2
add reference for biased estimators	3
uberleitungs blabla	3
outcome is possibly transformed: $h(T)$ instead of T	3
explain more detailed	4
add: counting process allows definition of (point) estimators, mathsabook p. 140 and Chapter 3.6	4
add parametric survival models	4
add reference for these tests	5
explain stratified cox model shortly OR add reference	5
allgemeines zu ML	6
add that mostly overfitting is a real problem + Quelle	6
add: difference of ML and DL	6
check benefits of doubly robust trafo in contrast to Buckley-James OR downsides of BJ	11
add sth about IPCW weights? add importance of positivity assumption (-> denominator)	12
add from basearticle: why IPCW is no longer considered/no suitable choice for estimator of full data loss	13
schauen ob Schreibweise konsistent ist, $zB \tilde{h}(T)$ etc	13
add reason why we need tau from Strawderman 2000	16
add paragraph that choice of h gives estimators for both, L2 and Brier loss	16
add names of algorithms for comparison	20
think of better title	21
change title	25
add reference Strawderman	25

Master Thesis

University Augsburg

Department of Mathematics

Chair for Computational Statistics
and Data Analysis



Principles of Deep Learning for Survival Analysis

Anna Sophie Rubeck

October 2021

Contents

1	Introduction	1
2	Principles of Survival Analysis	2
2.1	Censoring	2
2.2	Survival Analysis Data	3
2.3	Parametric and Semiparametric Approaches to Estimation of Survival Time	4
3	Principles of Deep Learning	6
3.1	Deep Feedforward Networks	6
4	Deep Learning for Fully Observed Data	8
5	Deep Learning for Censored Data	10
5.1	Censoring unbiased transformations	10
5.1.1	The Buckley-James Transformation	11
5.1.2	The Inverse Probability of Censoring Weighted Mapping	11
5.1.3	A Doubly Robust Censoring Unbiased Transformation	12
5.2	Censoring unbiased loss functions	13
5.3	Estimating Survival Probabilities $P(T \geq t X)$	15
5.4	Estimating Mean Survival	16
5.5	Implementation	16
6	Simulation of Survival Data	20
6.1	Simulation Settings	20
6.2	Implementation	21
6.2.1	Estimating the Survival Curves $S_0(. .)$ and $G_0(. .)$	21
6.2.2	Estimating m_k under L_2 -Loss	21
6.2.3	Selecting Truncation Time to Ensure Positivity	21
6.2.4	Algorithms used for comparison	21
6.3	Comparison	22
7	Conclusion	23
	Appendices	25
A	Proofs	25
A.1	Strawderman	25

1 Introduction

- growing importance of machine learning in many areas, as medicine - many different approaches - problem: treatment of missing data

2 Principles of Survival Analysis

This section is based on Kleinbaum and Klein (2020) and Klein and Moeschberger (2003).

In this section we will briefly describe the main ideas of survival analysis, the concept of censoring and parametric and semiparametric approaches to the problem of estimating survival times.

Survival analysis is used in various fields, as medicine, insurance and finance to predict a certain risk. The main goal is to estimate and analyze the time until a specified event occurs. An application for the medical field could for example be the time to death after the outbreak of a disease.

add more about the general setting (time series of data, consider medical cases as examples, ...)

In this thesis, we will consider one event of interest, that can either occur or not. It is also possible to consider competing risks. A detailed explanation of this case is for example given in Kalbfleisch and Prentice (2002, chapter 8). We also assume that the event of interest can only take place once. It is also possible to study recurrent events, i.e. events that can happen to a subject multiple times. For further details refer to Kalbfleisch and Prentice (2002, chapter 9).

2.1 Censoring

add reference

Survival data is often incomplete, which means that for some subjects the particular event is not observed within the study period. We refer to this data as censored.

Even though censored data is not observed until the event, it still contains valuable information. Assuming data arising from a medical study, there are three possible reasons why censoring occurs:

- The study ends before a participant experiences the event.
- A participant is lost to follow-up during the study period.
- A participant withdraws from the study or experiences a different event that makes further follow-up impossible.

The reasons mentioned above all lead to right censored data. We call the data right censored, if the true survival time is equal to or greater than the observed survival time.

There are two other types of censoring that may occur. If the survival time is less than or equal to the observed survival time, the data is left censored. This

could arise due to subjects that already experienced the event of interest before entering the study. A more general case is interval censoring. It occurs if we only know that the true survival time lies within specified time interval, but cannot observe it directly. The last case incorporates right censoring and left censoring as special cases.

Even though censored data is not fully observed, we cannot treat them the same as uncensored data or even leave them out completely, as this would both lead to biased estimators.

add reference for biased estimators

In this thesis we will consider the data to be possibly right censored, but not left or interval censored. Further information about left and interval censoring can for example be found in Kalbfleisch and Prentice (2002).

For further analysis, we make three assumptions about censoring, which are independent, random and non-informative censoring.

Random censoring

Independent censoring

non-informative censoring

These assumptions are covered more detailed in Kalbfleisch and Prentice (2002).

2.2 Survival Analysis Data

This section mainly relies on Kleinbaum and Klein (2020).

uberleitungs blabla

We denote by $T_i \geq 0$ the random variable for the survival time of subject i and by $C_i \geq 0$ the random variable for the censoring time of individual i . The observed time is given by $\tilde{T}_i = \min\{T_i, C_i\}$. The explanatory (or predictor) variables are given by the vector $X_i = (X_{i1}, \dots, X_{ip}) \in \delta \subset \mathbb{R}^p$, where δ denotes the bounded covariate space. The event indicator $d_i \in \{0, 1\}$ specifies whether subject i has experienced the event of interest ($d_i = 1$) or the data is censored ($d_i = 0$).

The basic data layout then has the form (\tilde{T}_i, d_i, X_i) , $i = 1, \dots, n$, where n is the number of patients in the study.

outcome is possibly transformed: h(T) instead of T

For analysis and prediction we often use two important functions: the survivor function and the hazard function. The survivor function $S(t) = P(T > t)$ specifies the distribution of the survival time. It can for example be used to analyze and compare the time to event of two different groups. The survivor function has three important properties: $S(t)$ is a non-increasing function, $S(0) = 1$ and $S(\infty) = 0$.

The hazard function $h(t)$, also called conditional failure rate, is given by

$$h(t) = \lim_{\Delta t \rightarrow 0} \frac{P(t \leq T < t + \Delta t | T \geq t)}{\Delta t}.$$

It is a measure of instantaneous potential, that can be used to identify a specific model form as exponential, Weibull or lognormal.

explain more detailed

The hazard function has two important properties: it is a non-negative function, i.e. $h(t) \geq 0 \forall t$ and it has no upper bound. In particular, this means that the hazard function can start anywhere in \mathbb{R}^+ and go up and down in any direction over time.

There exists a clearly defined relationship between the survivor function $S(t)$ and the hazard function $h(t)$:

$$S(t) = \exp \left(- \int_0^t h(u) du \right)$$

$$h(t) = - \left[\frac{dS(t)/dt}{S(t)} \right].$$

As a consequence, specifying one of the two function directly gives us the remaining.

- KM curves, log rank test

add: counting process allows definition of (point) estimators, mathsabook p. 140 and Chapter 3.6

2.3 Parametric and Semiparametric Approaches to Estimation of Survival Time

This section relies on Kleinbaum and Klein (2020).

add parametric survival models

The Cox Proportional Hazards Model (Cox PH model) is a popular approach to analyzing survival data. The main advantage is, that it is robust, i.e. it closely approximates the correct parametric model. In contrast to parametric models we do not need to assume the distribution of the outcome, which reduces the assumptions needed to properly fit the model.

The model is built on the following formulation of the hazard function:

$$h(t, x) = h_0(t) \exp \left(\sum_{i=1}^p \beta_i x_i \right),$$

where $x = (x_1, \dots, x_p)$ are the explanatory variables. This includes the central assumption of the Cox PH model, which is the proportional hazards (PH) assumption. It states that there exists a function $h_0(t)$, the baseline hazard, that does not involve the predictor variables and is thus independent of the observations.

As the baseline hazard $h_0(t)$ is not specified, the Cox PH model is a semiparametric model. Typically, maximum likelihood is used to estimate the parameters β_i , $i = 1, \dots, n$.

There are several tests to check, whether the PH assumption holds true, e.g. graphical techniques or goodness-of-fit-tests.

add reference for these tests

If this assumption is not fulfilled, we can use a stratified Cox procedure instead.

explain stratified cox model shortly OR add reference

In many cases it may be too simplistic to assume that the effect of the covariates is linear, as it is the case for the Cox PH model. Thus, we need a richer family of survival models to properly fit survival data with nonlinear risk functions.

One possible approach to this disadvantage is to use deep learning instead. In this thesis, we will take a closer look at this approach. But first, we will introduce the main ideas of deep learning and take a closer look at feed forward networks.

3 Principles of Deep Learning

This section relies on Goodfellow, Bengio, and Courville (2016).

allgemeines zu ML

Machine Learning algorithms consist of a task and a performance measure.

- classified according to what kind of experience they are allowed to have during the learning process

- learn/extract information from a given sample

- build prediction model according to the training sample

The central challenge for a machine learning algorithm is that it must perform well on new and previously unseen inputs, i.e. the error measure on the test set should be as small as possible and at the same time the error measure on the training set should also be reduced. This means that there are two factors that determine how well a machine learning algorithm will perform: make the training error small and make the gap between the training and the test error small.

If the first task is not fulfilled, underfitting occurs. This means that the error value on the training set is not sufficiently low for our model. If the second task is not fulfilled, overfitting occurs. In this case, the gap between training and test error is too large.

add that mostly overfitting is a real problem + Quelle

A common approach to reduce the possibility of overfitting is regularization. In principal, regularization is any modification of a learning algorithm, that aims to reduce its generalization error but not its training error. In practice this is often achieved by adding a penalty to the cost function.

add: difference of ML and DL

3.1 Deep Feedforward Networks

This section follows the ideas of Goodfellow, Bengio, and Courville (2016).

Deep feedforward networks are a commonly used deep learning model and aim to approximate some function f^* . More precisely, given the data x and the desired output y , the goal is to find f^* , such that $f^*(x) = y$. This is achieved by defining a mapping $y = f(x; \beta)$ and learning the parameters β that result in the best function approximation.

The term 'network' suggests that the mapping f is a composition of many different functions, i.e. $f(x) = f^{(K)} \circ \dots \circ f^{(1)}(x)$. The functions $f^{(i)}$ are called layers and the number of composed functions K gives the depth of the model.

The final layer is called the output layer. We want to achieve that this layer produces a value that is close to the desired output y , i.e. for each data x and

corresponding label y : $y \approx f^*(x)$.

The remaining layers are called hidden layers because we do not specify what exact output they generate. We are only interested in the prediction, the output of the final layer. The hidden layers are vector valued functions whose dimensions define the width of the model.

The networks are called feedforward networks because information only travels in one direction from input to output, i.e. there are no feedback connections between the layers.

Deep feedforward networks can extend linear models to overcome its limitations. The idea is to apply the linear model not directly to the data x but to a transformed input $\phi(x)$, where ϕ is a nonlinear transformation. The remaining question is, how to choose this mapping ϕ . The deep learning approach to this problem is to define the mapping $y = \phi(x, \beta)^T w$ and learn the parameter β . This is an example for a deep feedforward network with one hidden layer, namely ϕ .

soil rein:

- model is associated with a directed acyclic graph
- most loss functions non convex for neural networks -> use iterative, gradient based optimizers. Important: initialize all weights to small random values
- important aspect of designing deep feedforward model: choice of cost function
- most neural networks are trained using maximum likelihood -> cost function = negative log likelihood (gradient of cost function must be large and predictable enough to serve as a good guide for the learning algorithm -> use sigmoid output unit(for binomial), use softmax for multinoulli); most popular cost function: cross entropy cost function
- when using feedforward networks: need to determine how many units we need and how they are connected

4 Deep Learning for Fully Observed Data

This section follows the ideas of Steingrímsson and Morrison (2020) and insights described in Goodfellow, Bengio, and Courville (2016).

As described in section 2.2, we assume that the dataset consists for each subject $i = 1, \dots, n$ of a positive continuous failure time $T_i \in \mathbb{R}^+$ and a covariate vector $X_i = (X_{i1}, \dots, X_{ip}) \in \delta \subset \mathbb{R}^p$. The set δ describes the bounded covariate space. Again, the outcome T_i is possibly transformed via a monotone function $h : \mathbb{R}^+ \rightarrow \mathbb{R}$. Possible choices for this transformation $h(T_i)$ are the identity function $h(T_i) = T_i$ or the logarithmic transformation $h(T_i) = \ln(T_i)$.

In this section we assume that the data is not censored. In this case our data set reduces to the fully observed data given by

$$\mathcal{F} = \{(h(T_i), X_i); i = 1, \dots, n\}.$$

We will now describe the structure of a deep learning algorithm based on feedforward networks for fully observed data, as presented in Steingrímsson and Morrison (2020).

Algorithm 4.1.

1. Setup the structure of the risk prediction model.
2. Estimate the weight vector.
3. Use cross-validation to select the penalization parameter η from a predetermined sequence η_1, \dots, η_M .
4. Create the final prediction model.

We will now take a closer look at these steps.

First we define the layout of the neural network we want to use for our prediction. For this, we need to set the number of layers K and define the functions $f_{\beta_k}^{(k)}(X)$ for each layer $k = 1, \dots, K$. The output of the hidden layer architecture is then given by $f_{\beta}(X) = f_{\beta_K}^{(K)} \circ f_{\beta_{K-1}}^{(K-1)} \circ \dots \circ f_{\beta_1}^{(1)}(X)$, where $\beta = (\beta_1^T, \dots, \beta_K^T)^T$ is a vector of unknown weights which will be estimated next.

In the second step, the estimation of the weight vector, we need a prespecified loss function $L(h(T), f(X))$, that describes the difference between the prediction of our model $f(X)$ and the desired outcome $h(T)$. We also use a penalization parameter η , that tries to keep the weights as small as possible. Large weights can be a sign of a complex network and thus overfitting of the training data. To counteract this effect, we penalize such large weights.

Now we want to estimate the vector of weights β by minimizing the empirical penalized loss function $\hat{L}(\beta)$, which has the following form:

$$\hat{L}(\beta) = \frac{1}{n} \sum_{i=1}^n L(h(T_i), f_{\beta}(X_i)) + \eta \|\beta\|^2, \quad (4.1)$$

where $\|\beta\|^2$ is the L_2 norm of β , given by $\|\beta\|^2 = \sum_{i=1}^q |\beta_i|^2$.

The third step ensures, that we make a suitable choice for the penalization parameter η . To perform cross-validation, we first split the given dataset into D disjoint sets K_1, \dots, K_D . Then, for fixed $l \in \{1, \dots, D\}$ and $m \in \{1, \dots, M\}$, we set $\hat{\beta}_{\eta_m}^{(l)}$ as the vector of weights estimated in equation (4.1) using the penalization parameter η_m and the data \mathcal{F}/K_l . Let $\delta_{i,l}$ be an indicator whether observation i is included in the dataset K_l ($\delta_{i,l} = 1$) or not ($\delta_{i,l} = 0$). The cross-validation error corresponding to the penalization parameter η_m is then defined as

$$\alpha(m) = \sum_{l=1}^D \sum_{i=1}^n \delta_{i,l} L(h(T_i), f_{\hat{\beta}_{\eta_m}^{(l)}}(X_i)).$$

We then choose the penalization parameter, that minimizes the cross-validation error, i.e. let $\eta = \eta_{m^*}$, where $m^* = \arg \min_{m \in \{1, \dots, M\}} \alpha(m)$.

In step four we use the optimal penalization parameter η_{m^*} to estimate the vector of weights $\hat{\beta}_{\eta_{m^*}}^{(l)}$ and then build the final prediction model $f_{\hat{\beta}_{\eta_{m^*}}^{(l)}}(X)$.

The estimated vector of weights β minimizes the expected loss used in the algorithm. Thus, the parameter that we want to estimate determines the choice of loss function. If we use for example the L_2 loss, the population parameter estimated by the full data deep learning algorithm is the conditional mean $E[h(T)|X]$. If we instead want to estimate survival probabilities $P(T \geq t|X)$ for a fixed time point t , we can use the Brier loss which is given by $E[(I(T \geq t) - \beta(X))^2]$.

5 Deep Learning for Censored Data

This section relies on Steingrimsdóttir and Morrison (2020) as well as Goodfellow, Bengio, and Courville (2016).

In this section we will now allow for the dataset to contain censored data, which means that in some cases we cannot observe the true failure time of an individual. The main difficulty in extending algorithm 4.1 to possibly censored data is to define a suitable loss function, that can be calculated in the presence of censoring and at the same time reduces to the full data loss $L(h(T), f(X))$ if the data is not censored.

The observed dataset consists of n independent and identically distributed observations and is given by

$$\mathcal{O} = \{(\tilde{T}_i = \min\{T_i, C_i\}, \Delta_i = I(T_i \leq C_i), X_i); i = 1, \dots, n\},$$

where \tilde{T}_i denotes the observed time and Δ_i indicates, whether subject i experienced the event of interest or was censored.

We define the conditional survivor functions for T and C as $S_0(u|X) = P(T > u|X)$ and $G_0(u|X) = P(C > u|X)$ respectively. We assume that the censoring time C is continuous and that the censoring mechanism is non-informative, i.e. that C is independent of $T|X$. For further calculations we also assume that $G_0(\tilde{T}|X) \geq \varepsilon > 0$ for some $\varepsilon > 0$, to ensure positivity of the conditional censoring distribution.

To close the gap between what is done in the presence and absence of censoring, we will now define censoring unbiased transformations.

5.1 Censoring unbiased transformations

This section is based on Steingrimsdóttir, Diao, and Strawderman (2019) and Rubin and van der Laan (2007).

For prediction with right censored data it is a popular approach to replace the possibly censored responses with surrogate values. This can be achieved by using an appropriate mapping $Y^*(\cdot)$. The transformed values are then entered in standard smoothing algorithms.

In our case the key requirement is, that the mapping $Y^*(\cdot)$ is a censoring unbiased transformation, which is defined as follows.

Definition 5.1. Let Y be a scalar function of the full data \mathcal{F} and $Y^*(\mathcal{O})$ a scalar function of the observed data \mathcal{O} .

Then we call $Y^*(\mathcal{O})$ a *censoring unbiased transformation (CUT)* for Y if for every point x in the covariate space δ , i.e. $x \in \delta \subset \mathbb{R}^p$:

$$E[Y^*(\mathcal{O})|X = x] = E[Y|X = x].$$

In other words, a censoring unbiased transformation maps the response in a way that it keeps its conditional mean structure.

We will now take a closer look at three specific CUTs: the Buckley-James Transformation, the inverse probability of censoring weighted mapping and a doubly robust censoring unbiased transformation.

5.1.1 The Buckley-James Transformation

Rubin and van der Laan (2007)

The Buckley-James Transformation is given by

$$Y_{BJ}^*(\mathcal{O}) = \Delta Y + (1 - \Delta)m(Y, X; S), \quad (5.1)$$

where

$$m(t, x; S) = E_S[T|T > t, X = x] = \frac{1}{S(t|X = x)} \int_t^\infty u dS(u|X = x) \quad (5.2)$$

and $S(\cdot|X)$ denotes the conditional survival function.

This Transformation is rather intuitive. For an uncensored subject, i.e. $\Delta = 1$, the response remains unchanged. If the response is censored, i.e. $\Delta = 0$, it is mapped to its conditional expectation.

The Buckley-James Transformation has an important property. Among all CUTs $Y^*(\mathcal{O})$ the mapping $Y_{BJ}^*(\mathcal{O})$ results in the best predictor of the original response, i.e.

$$Y_{BJ}^*(\mathcal{O}) = \arg \min_{Y^*(\mathcal{O}) \text{ is a CUT}} E|Y^*(\mathcal{O}) - Y|^2.$$

On the other hand, it requires the correct specification of the survivor function $S(\cdot|X)$.

check benefits of doubly robust trafo in contrast to Buckley-James OR down-sides of BJ

5.1.2 The Inverse Probability of Censoring Weighted Mapping

Rubin and van der Laan (2007) and Steingrimsson, Diao, Molinaro, and Strawderman (2016).

In contrast to the Buckley-James transformation, which depends on the survivor function $S(\cdot|X)$, it is also possible to build mappings that instead depend only on the censoring mechanism. One example for this is the inverse probability of censoring weighted (IPCW) mapping. Here, the basic idea is to weight the

contribution of each uncensored observation by the inverse probability of being censored. The IPCW mapping is defined as follows:

$$Y_{IPCW}^*(\mathcal{O}) = \frac{\Delta Y}{G(Y|X)}.$$

When there is no censored data in the whole data set, the response remains unchanged.

For the IPCW transformation we need to estimate the conditional distribution of the censoring times given the covariates $G(\cdot|X)$.

add sth about IPCW weights? add importance of positivity assumption (-> denominator)

5.1.3 A Doubly Robust Censoring Unbiased Transformation

Rubin and van der Laan (2007)

The Buckley-James estimator, as well as the IPCW mapping, depend on the nuisance parameters $S(\cdot|X)$ and $G(\cdot|X)$ respectively. As a consequence, poor estimators of these two conditional distributions can significantly degrade the performance of regression applied to the transformed data.

A third CUT, that requires a well approximation of either $S(\cdot|X)$ or $G(\cdot|X)$ (but not both), is the so called doubly robust CUT. It is defined as follows:

$$Y_{DR}^*(\mathcal{O}) = \frac{\Delta Y}{G(Y|X)} + \frac{(1 - \Delta)m_S(Y, X)}{G(Y|X)} - \int_0^Y \frac{m(c, X; S)}{G^2(c|X)} dG(c|X), \quad (5.3)$$

with $m(c, X; S)$ as defined in equation (5.2).

The first term in equation (5.3) is equal to the IPCW transformation. Similar to the Buckley-James transformation, a term that considers the censored data as well is added. The last part can be seen as an augmentation term. In total, the doubly robust transformation can be seen as an augmented inverse probability weighted mapping. For further details about the augmentation refer to Kalbfleisch and Prentice (2002).

The double robustness property is formalized in theorem 1 in Rubin and van der Laan (2007) and is reproduced below.

Theorem 5.1. *Suppose that the response Y and the censoring time C are continuous random variables, that they are conditionally independent given the covariates X , and that the conditional distribution $G_0(C|X)$ has a conditional density $g_0(\cdot|X)$. Assume that $Y \leq \tau < \infty$ for some τ and that $S(\tau|X) = 0$ with probability one for some conditional function $\hat{S}(\tau|X)$. Suppose further that $\tilde{G}(\tau|X) \geq \varepsilon > 0$ for*

some ε and conditional survival function $\tilde{G}(\tau|X)$, with corresponding conditional density $\tilde{g}(\cdot|X)$. Assume that $\tilde{g}(\cdot|X = x)$ is absolutely continuous with respect to $g_0(\cdot|X = x)$ for all x . We will use the convention that $m(y, x; \tilde{S})$ is set to zero if $\tilde{S}(y|X = x) = 0$. Then, if either $S_0(\cdot|X)$ or $G_0(\cdot|X)$ are specified correctly, i.e. $S_0(\cdot|X) = \tilde{S}(\cdot|X)$ or $G_0(\cdot|X) = \tilde{G}(\cdot|X)$, it holds that

$$E[Y_{DR}^*(\mathcal{O})|X] = E[Y|X].$$

5.2 Censoring unbiased loss functions

Steingrimsdóttir, Diao, and Strawderman (2019) and Steingrimsdóttir and Morrison (2020).

The theory of censoring unbiased transformations can be directly applied to loss functions and results in a class of censoring unbiased loss functions, that close the gap between what is done in the presence and absence of censoring.

We will focus on two censoring unbiased estimators for the estimated loss $\mathcal{R}(\beta) = E[L(h(T), \beta(X))]$ that both reduce to the full data loss in the absence of censoring: the doubly robust loss function and the Buckley-James estimator.

add from basearticle: why IPCW is no longer considered/no suitable choice for estimator of full data loss

The Buckley-James estimator for $\mathcal{R}(\beta)$ is given by

$$L_{BJ}(\mathcal{O}, \beta; S_0) = \frac{1}{n} \sum_{i=1}^n [\Delta_i L(h(T_i), \beta(X_i)) + (1 - \Delta_i) m_L(C_i, X_i, \beta; S_0)]. \quad (5.4)$$

This estimated loss function is the average of the Buckley-James transformed full data loss.

Again, the structure of this estimator is rather intuitive. For uncensored observations we add the full data loss $L(h(T_i), \beta(X_i))$ and for censored observations we use the expectation of the full data loss $m_L(C_i, X_i, \beta; S_0)$ given by equation (5.6).

In order to consistently estimate $\mathcal{R}(\beta)$, this loss function requires a consistent estimator for the conditional survival function $S_0(\cdot|X)$.

The doubly robust loss function is defined as

schauen ob Schreibweise konsistent ist, zB tilde(T) etc

$$\begin{aligned}
L_{DR}(\mathcal{O}, \beta; G_0, S_0) &= \frac{1}{n} \sum_{i=1}^n \frac{\Delta_i L(h(T_i), \beta(X_i))}{G_0(\tilde{T}_i|X_i)} \\
&+ \frac{1}{n} \sum_{i=1}^n \left(\frac{(1 - \Delta_i) m_L(\tilde{T}_i, X_i, \beta; S_0)}{G_0(\tilde{T}_i|X_i)} \right. \\
&\quad \left. - \int_0^{\tilde{T}_i} \frac{m_L(u, X_i, \beta; S_0)}{G_0(u|X_i)} d\Lambda_{G_0}(u|X_i) \right), \tag{5.5}
\end{aligned}$$

where for a survival curve S

$$\begin{aligned}
m_l(u, x, \beta; S) &= E_S[L(h(T), \beta(X)) | T \geq u, X = x] \\
&= - \int_u^\infty \frac{L(h(t), \beta(x))}{S(u|w)} dS(t|w) \tag{5.6}
\end{aligned}$$

and $\Lambda_G(u|X) = - \int_0^u 1/G(t|X) dG(t|X)$ is the cumulative hazard function.

Again, we obtain this loss function by calculating the average of the full data loss that is now transformed via equation (5.3).

Theorem 5.1 also holds for the doubly robust loss function. As a consequence we only need to model either $S(\cdot|X)$ or $G(\cdot|X)$ correctly, to obtain a suitable estimator.

As a next step we can plug in estimators \hat{S} and \hat{G} for the conditional survival functions S_0 and G_0 respectively to obtain empirical estimators for the full data loss.

For the empirical Buckley-James loss we obtain

$$L_{BJ}(\mathcal{O}, \beta; \hat{S}) = \frac{1}{n} \sum_{i=1}^n \left(\Delta_i L(h(T_i), \beta(X_i)) + (1 - \Delta_i) m_L(C_i, X_i, \beta; \hat{S}) \right).$$

The empirical doubly robust loss function is given by

$$\begin{aligned}
L_{DR}(\mathcal{O}, \beta; \tilde{G}, \tilde{S}) &= \frac{1}{n} \sum_{i=1}^n \frac{\Delta_i L(h(T_i), \beta(X_i))}{\tilde{G}(\tilde{T}_i|X_i)} \\
&+ \frac{1}{n} \sum_{i=1}^n \left(\frac{(1 - \Delta_i) m_L(\tilde{T}_i, X_i, \beta; \tilde{S})}{\tilde{G}(\tilde{T}_i|X_i)} \right. \\
&\quad \left. - \int_0^{\tilde{T}_i} \frac{m_L(u, X_i, \beta; \tilde{S})}{\tilde{G}(u|X_i)} d\Lambda_{\tilde{G}}(u|X_i) \right).
\end{aligned}$$

If we use the incorrect model specification, that $\hat{G}(t, x) = 1$ for all pairs (t, x) , the empirical Buckley-James loss is equivalent to the empirical doubly robust loss, i.e. $L_{BJ}(\mathcal{O}, \beta; \hat{S}) = L_{DR}(\mathcal{O}, \beta; \hat{S}, \hat{G} = 1)$.

We can now replace the full data loss function $L(h(T), f(X))$ in algorithm 4.1 with either the Buckley-James loss L_{BJ} or the doubly robust loss L_{DR} and obtain a prediction model, that can be calculated with the observed data \mathcal{O} . We refer to the resulting algorithms as the Buckley-James or the doubly robust deep learning algorithms respectively. As both loss functions are censoring unbiased loss functions, we refer to the resulting algorithms as censoring unbiased deep learning (CUDL).

5.3 Estimating Survival Probabilities $P(T \geq t|X)$

This section is based on Steingrimssohn and Morrison (2020).

In this section we will specify how we can estimate the survival probabilities $P(T \geq t|X)$ for a specified time point t . As mentioned before, the target parameter determines the choice of loss function. For the estimation of survival probabilities we can use the Brier loss $L_{t,2}(T, \beta(X)) = (I(T \geq t) - \beta(X))^2$ as our full data loss.

As a first step we have to define a modified dataset with respect to the fixed time point t :

$$\mathcal{O}(t) = \{(\tilde{T}_i(t) = \min(T_i, C_i, t), \Delta_i(t) = I(\min(T_i, t) \leq C_i), X_i); i = 1, \dots, n\}.$$

Using the Brier loss as the full data loss in equation (5.5) results the doubly robust Brier loss, which is given by

$$\begin{aligned} L_{DR,t}(\mathcal{O}(t), \beta; \hat{G}, \hat{S}) &= \frac{1}{n} \sum_{i=1}^n \frac{\Delta_i(t)(I(\tilde{T}_i \geq t) - \beta(X_i))^2}{\hat{G}(\tilde{T}_i|X_i)} \\ &\quad + \frac{1}{n} \sum_{i=1}^n \frac{(1 - \Delta_i(t))m_{L_2,t}(\tilde{T}_i(t), X_i, \beta; \hat{S})}{\hat{G}(\tilde{T}_i(t)|X_i)} \\ &\quad - \int_0^{\tilde{T}_i(t)} \frac{m_{L_2,t}(u, X_i, \beta; \hat{S})d\hat{\Lambda}_G(u|X_i)}{\hat{G}(u|X_i)}, \end{aligned}$$

with

$$m_{L_2,t}(u, X, \beta; S) = E_S[(I(\tilde{T} \geq t) - \beta(X))^2 | T \geq u, X].$$

Analogously, by plugging in the Brier risk to the Buckley-James estimators we obtain the empirical Buckley-James Brier loss given by

$$L_{BJ,t}(\mathcal{O}, \beta; \hat{S}) = \frac{1}{n} \sum_{i=1}^n \left(\Delta_i(t)(I(\tilde{T}_i \geq t) - \beta(X_i))^2 + (1 - \Delta_i(t))m_{L_2,t}(\tilde{T}_i(t), X_i; \hat{S}) \right)$$

Both loss functions can now be incorporated in the CUDL algorithm and result in a prediction model for $P(T \geq t|X)$.

5.4 Estimating Mean Survival

This section follows the ideas presented in Steingrimsdottir and Morrison (2020) and Strawderman (2000).

An other popular goal is to estimate the mean survival $E[T|X]$. For this target parameter we use the L_2 loss as our full data loss.

add reason why we need tau from Strawderman 2000

This is why we alternatively estimate the restricted mean survival $E[\min(T, \tau)|X]$ for a fixed constant τ .

Analogously to the estimation of survival probabilities, we need to incorporate the constant τ in the data set. This leads us to the modified version

$$\mathcal{O}(\tau) = \{\tilde{T}_i(\tau) = \min\{T_i, C_i, \tau\}, \Delta_i(\tau) = I(\min\{T_i, \tau\} \leq C_i), X_i; i = 1, \dots, n\}.$$

Selecting the L_2 loss as the full data loss and applying the CUDL algorithms to the modified data set gives us an estimator for the restricted mean survival.

5.5 Implementation

This section relies on Steingrimsdottir and Morrison (2020).

We will now describe how we can implement the doubly robust and Buckley-James CUDL algorithm using a specific response transformation.

add paragraph that choice of h gives estimators for both, L2 and Brier loss

We will use the following response transformation:

$$D(O_i; G, S) = A_{1i}(G) + B_{1i}(G, S) - C_{1i}(G, S),$$

where for $k = 0, 1, 2$

$$A_{ki} = \frac{\Delta_i h(\tilde{T}_i)^k}{G(\tilde{T}_i|X_i)},$$

$$B_{ki} = \frac{(1 - \Delta_i) m_k(\tilde{T}_i, X_i; S)}{G(\tilde{T}_i|X_i)}$$

and

$$C_{ki}(G, S) = \int_0^{\tilde{T}_i} \frac{m_k(u, X_i; S)}{G(u|X_i)} d\Lambda_G(u|X_i).$$

For $k = 1, 2$

$$m_k(t, x; S) = E_S[h^k(T)|T \geq t, X = x] = -[S(t|x)]^{-1} \int_t^\infty [h(u)^k] dS(u|x)$$

and $m_0(t, x; S) = 1$ for all pairs (t, x) .

The transformation described above is equal to the doubly robust CUT described in section 5.1.3 and thus requires that at least one of the models $T|X$ or $C|X$ is correctly specified.

We can now define the response transformed L_2 loss, that uses the censoring unbiased outcome transformation $D(O_i; G, S)$ as the outcome:

$$L_2^*(\mathcal{O}, \beta; G, S) = \frac{1}{n} \sum_{i=1}^n (D(O_i; G, S) - \beta(X_i))^2.$$

Our main goal is to show that a deep learning algorithm, that uses the transformed data and the response transformed L_2 loss, is equivalent to the CUDL algorithms presented in section 5.2. As a result, the CUDL algorithms inherits the desired doubly robustness property to the transformed data algorithm.

This equivalence is stated in theorem 1 in Steingrimssohn and Morrison (2020) and reproduced below.

Theorem 5.2. *Under mild regularity conditions, the prediction model created using the CUDL algorithm with the loss function $L_{DR}^{(2)}(\mathcal{O}, \beta; G, S)$ is identical to the prediction model built using the full data deep learning algorithm implemented using the loss function $L_2^*(\mathcal{O}, \beta; G, S)$.*

Proof. Using the notation we introduced earlier in this section, we can rewrite the doubly robust loss function with the L_2 loss as the full data loss as:

$$\begin{aligned} L_{DR}^{(2)}(\mathcal{O}, \beta; G, S) &= \frac{1}{n} \sum_{i=1}^n ((A_{2i} + B_{2i} - C_{2i}) - 2(A_{1i} + B_{1i} - C_{1i})\beta(X_i) \\ &\quad + (A_{0i} + B_{0i} - C_{0i})\beta(X_i)^2). \end{aligned} \quad (5.7)$$

This equality can be shown by reordering the terms in the sum above in alphabetical order. Lemma A.1 provided in the appendix states that $A_{0i} + B_{0i} - C_{0i} = 1$ holds true. This leads us to

$$L_{DR}^{(2)}(\mathcal{O}, \beta; G, S) = \frac{1}{n} \sum_{i=1}^n ((A_{2i} + B_{2i} - C_{2i}) - 2(A_{1i} + B_{1i} - C_{1i})\beta(X_i) + \beta(X_i)^2).$$

We can also reformulate the response transformed L_2 loss by calculating the square and obtain

$$L_2^*(\mathcal{O}, \beta; G, S) = \frac{1}{n} \sum_{i=1}^n ((A_{1i} + B_{1i} - C_{1i})^2 - 2(A_{1i} + B_{1i} - C_{1i})\beta(X_i) + \beta(X_i)^2). \quad (5.8)$$

We can see that the loss functions in equation (5.7) and equation (5.8) are actually very similar. Only the first term in the sum, which is independent of β , is different for each loss function.

As a consequence for a fixed penalization parameter η minimizing

$$L_2^*(\mathcal{O}, \beta; G, S) + \eta \|\beta\|_p^2$$

and

$$L_{DR}^{(2)}(\mathcal{O}, \beta; G, S) + \eta \|\beta\|_p^2$$

results in the same weight vector β .

If we can show that the penalization parameter selected by minimizing the cross-validated loss with the loss $L_2^*(\mathcal{O}, \beta; G, S)$ is the same if we replace the L_2^* loss by $L_{DR}^{(2)}(\mathcal{O}, \beta; G, S)$, both prediction models produce the identical output.

To show this, we use the notation from section XXX. Recall that $\delta_{i,l}$ indicates whether observation i is an element of the subset K_l .

As $L_2^*(\mathcal{O}, \beta; G, S)$ and $L_{DR}^{(2)}(\mathcal{O}, \beta; G, S)$ are equal up to a term that does not depend on β , it holds that

$$\sum_{l=1}^D \sum_{i=1}^n \delta_{i,l} L_2^*(\mathcal{O}, \beta; G, S) = \sum_{l=1}^D \sum_{i=1}^n \delta_{i,l} L_{DR}^{(2)}(\mathcal{O}, \beta; G, S) + \text{const}(\mathcal{O}; G, S).$$

This shows that

$$\begin{aligned} & \arg \min_{m \in \{1, \dots, M\}} \sum_{l=1}^D \sum_{i=1}^n \delta_{i,l} L_2^*(\mathcal{O}, \beta; G, S) \\ &= \arg \min_{m \in \{1, \dots, M\}} \sum_{l=1}^D \sum_{i=1}^n \delta_{i,l} L_{DR}^{(2)}(\mathcal{O}, \beta; G, S), \end{aligned}$$

which completes the proof. \square

The assumptions on the conditional censoring distribution $G(\cdot|x)$ are flexible enough to allow for $G(t|x) = 1$ for all pairs (t, x) . And as it holds that $L_{BJ}(\mathcal{O}, \beta; \hat{S}) = L_{DR}(\mathcal{O}, \beta; \hat{S}, \hat{G} = 1)$, theorem 5.2 also applies to the CUDL using the Buckley-James loss.

Theorem 5.2 allows us to implement the CUDL algorithm using software for fully observed outcomes. This leads us to the following algorithm, as described in Steingrímsson and Morrison (2020):

Algorithm 5.1.

1. Use the observed data \mathcal{O} to calculate estimators for the survival curves $\hat{S}(.|.)$ and $\hat{G}(.|.)$
2. Perform the response transformation $D(O_i; \hat{G}, \hat{S}), i = 1, \dots, n$
3. Run software implementing the L_2 full data deep learning algorithm on the dataset $\{(D(O_i; \hat{G}, \hat{S}), X_i); i = 1, \dots, n\}$.

6 Simulation of Survival Data

In this section we compare the performance of XXX

add names of algorithms for comparison

6.1 Simulation Settings

The different settings for our simulation follow those presented in Steingrimsdottir, Diao, and Strawderman (2019).

Setting 1. In our first setting we simulate the dataset such that the proportional hazards assumption is met. To achieve this we create 300 independent observations with 20 covariates $W = (W_1, \dots, W_{20})$. The covariate vector is multivariate normal with mean zero and a covariance matrix M with elements $m_{ij} = 0.9^{|i-j|}$. The survival times are simulated from an exponential distribution with mean $\mu = \exp(0.1 \sum_{i=9}^{18} W_i)$. The censoring distribution is also simulated from an exponential distribution with mean XXX.

Setting 2. In our second setting the proportional hazards assumption is mildly violated. Again, we simulate 300 independent observations with 20 covariates $W = (W_1, \dots, W_{20})$. The covariate vector now consists of 20 iid uniform random variables on the interval $[0, 1]$. The survival time follows an exponential distribution with mean $\mu = \sin(W_1\pi) + 2|W_2 - 0.5| + W_3^3$. The censoring times are uniformly distributed on the interval $[0, 6]$. This results in approximately XXX% censoring.

Setting 3. In this setting, the proportional hazards assumption is strongly violated. Our dataset consists of 300 independent observations and the covariate vectors $W = (W_1, \dots, W_{20})$ are multivariate normal with mean zero and covariance matrix M with $m_{ij} = 0.75^{|i-j|}$. Now, the survival times are simulated using the gamma distribution with shape parameter $\mu = 0.5 + 0.3|\sum_{i=11}^{15} W_i|$ and scale parameter $\lambda = 2$. The censoring times are uniformly distributed on the interval $[0, 15]$. This leads to approximately XXX% censoring.

Setting 4. In the last setting we consider the case where the censoring distribution depends on covariates. Again, we simulate 300 independent observations where the covariates $W = (W_1, \dots, W_{20})$ are multivariate normal with mean zero and the covariance matrix M with $m_{ij} = 0.75^{|i-j|}$. The survival distribution is now a log-normal distribution with mean $\mu = 0.1|\sum_{i=1}^5 W_i| + 0.1|\sum_{i=16}^{20} W_i|$. The censoring times are also simulated according to a log-normal distribution with mean $\tilde{\mu} = \mu + 0.5$ and scale parameter $\lambda = 1$. In this setting approximately XXX% of our data are censored.

6.2 Implementation

which algorithms are used

6.2.1 Estimating the Survival Curves $S_0(.|.)$ and $G_0(.|.)$

This section follows the ideas presented in Rubin and van der Laan (2007) and Steingrimsdottir, Diao, Molinaro, and Strawderman (2016).

If the censoring time is completely independent of both, the event time T and the covariates X , the conditional survival function $\tilde{G}(.|X) = \tilde{G}(.)$ can be efficiently estimated with the Kaplan-Meier estimator. For example if we consider that censoring is caused by the end of the study, the above stated independence assumption is met.

To minimize the required number of assumptions, we chose to estimate the censoring survival curve $G_0(t|X)$ using the survival tree method. This tree based method extends the ph regression to tree-structured relative risk functions. Further details can be found in LeBlanc and Crowley (1992).

The survival curve $S_0(t|X)$ should not be estimated using the Kaplan-Meier estimator, as we presume that the observed time T depends on the covariates X . Instead, the survival curve is estimated via a random survival forest procedure. For a detailed explanation refer to Steingrimsdottir, Diao, Molinaro, and Strawderman (2016).

It is possible to choose different estimators for $S_0(.|X)$ and $G_0(.|X)$ respectively.

When choosing the estimators it is important to avoid using an estimator for one of the functions that relies on the estimator of the other. This is because it would negatively impact the doubly robustness property, if we specify one of the estimators incorrectly and thus affect the consistency of the other. Further details can be found in Steingrimsdottir, Diao, Molinaro, and Strawderman (2016).

6.2.2 Estimating m_k under L_2 -Loss

This section is based on Steingrimsdottir, Diao, Molinaro, and Strawderman (2016).

6.2.3 Selecting Truncation Time to Ensure Positivity

Steingrimsdottir, Diao, Molinaro, and Strawderman (2016).

6.2.4 Algorithms used for comparison

think of better title

Cox and penalized cox model
survival trees

6.3 Comparison

We will now compare the results using different methods for time to event estimation.

7 Conclusion

evtl als Ausblick: Erweiterung auf competing risks, time dependent covariates, recurrent events; Test different estimators for S and G

References

- Goodfellow, I., Y. Bengio, and A. Courville (2016). *Deep Learning* (first ed.). The MIT Press.
- Kalbfleisch, J. D. and R. L. Prentice (2002). *The statistical Analysis of Failure Time Data* (second ed.). John Wiley & Sons.
- Klein, J. P. and M. L. Moeschberger (2003). *Survival Analysis: Techniques for Censored and Truncated Data* (second ed.). Springer.
- Kleinbaum, D. G. and M. Klein (2020). *Survival Analysis* (third ed.). Springer Science + Business Media.
- LeBlanc, M. and J. Crowley (1992). Relative risk trees for censored survival data. *International Biometric Society Vol. 48, No. 2*, 411–425.
- Rubin, D. and M. J. van der Laan (2007). A doubly robust censoring unbiased transformation. *The International Journal of Biostatistics Vol. 3, Iss. 1*, Article 4.
- Steingrimsson, J. A., L. Diao, A. M. Molinaro, and R. L. Strawderman (2016). Doubly robust survival trees. *Statistics in medicine* 35.
- Steingrimsson, J. A., L. Diao, and R. L. Strawderman (2019). Censoring unbiased regression trees and ensembles. *Journal of the American Statistical Association* 114,525, 370–383.
- Steingrimsson, J. A. and S. Morrison (2020). Deep learning for survival outcomes. *Statistics in medicine* 39 (17), 2339–2349.
- Strawderman, R. L. (2000). Estimating the mean of an increasing stochastic process at a censored stopping time. *Journal of the American Statistical Association* 95, No. 452, 1192–1208.

Appendices

A Proofs

A.1 Strawderman

change title

Lemma A.1. *Using the notation introduced in section XXX, it holds that*

$$A_{0i} + B_{0i} - C_{0i} = 1$$

Proof. By plugging in the definitions of A_{0i} , B_{0i} and C_{0i} respectively, we obtain

$$A_{0i} + B_{0i} - C_{0i} = \frac{\Delta_i}{G(T_i|X_i)} + \frac{1 - \Delta_i}{G(T_i|X_i)} - \int_0^{T_i} \frac{1}{G(u|X_i)} d\Lambda_G(u|X_i)$$

We can now combine the first two terms to obtain

$$A_{0i} + B_{0i} - C_{0i} = \frac{1}{G(T_i|X_i)} - \int_0^{T_i} \frac{1}{G(u|X_i)} d\Lambda_G(u|X_i). \quad (\text{A.1})$$

This formulation allows us to apply lemma 1 from Strawderman (2000). Using the notation introduced in this thesis, the aforementioned lemma gives us the following equation:

$$\frac{\Delta_i^*}{G(T_i|X_i)} = 1 - \int_0^\infty \frac{1}{G(s|X_i)} dM^*(s), \quad (\text{A.2})$$

with

$$M^*(s) = I\{T_i \leq s, \Delta_i^* = 0\} - \int_0^s I\{T_i \geq u\} d\Lambda_G(u|X_i),$$

and Δ^* is any arbitrary indicator function.

Now we set $\Delta_i^* = 1$ for all $i = 1, \dots, n$ and take a closer look at the definition of the function $M^*(s)$.

$$\begin{aligned} M^*(s) &= I\{T_i \leq s, \Delta_i^* = 0\} - \int_0^s I\{T_i \geq u\} d\Lambda_G(u|X_i) \\ &= - \int_0^s I\{T_i \geq u\} d\Lambda_G(u|X_i) \\ &= - \int_0^{\min\{T_i, s\}} d\Lambda_G(u|X_i) \\ &= - (\Lambda_G(\min\{T_i, s\}|X_i) - \Lambda_G(0|X_i)) \\ &= -\Lambda_G(\min\{T_i, s\}|X_i) \end{aligned}$$

Now we can plug in our choice of Δ_i^* and $M^*(s)$ to equation (A.2).

$$\begin{aligned}
\frac{1}{G(T_i|X_i)} &= 1 + \int_0^\infty \frac{1}{G(s|X_i)} d\Lambda_G(\min\{T_i, s\}|X_i) \\
&= 1 + \int_0^{T_i} \frac{1}{G(s|X_i)} d\Lambda_G(s|X_i) + \int_{T_i}^\infty \frac{1}{G(s|X_i)} d\Lambda_G(T_i|X_i) \\
&= 1 + \int_0^{T_i} \frac{1}{G(s|X_i)} d\Lambda_G(s|X_i)
\end{aligned} \tag{A.3}$$

Combining equation (A.1) and (A.3) now yields:

$$\begin{aligned}
A_{0i} + B_{0i} - C_{0i} &= \frac{1}{G(T_i|X_i)} - \int_0^{T_i} \frac{1}{G(u|X_i)} d\Lambda_G(u|X_i) \\
&= 1 + \int_0^{T_i} \frac{1}{G(s|X_i)} d\Lambda_G(s|X_i) - \int_0^{T_i} \frac{1}{G(u|X_i)} d\Lambda_G(u|X_i) \\
&= 1.
\end{aligned}$$

□